

IE7275 12159 Data Mining in Engineering SEC 02 Fall 2024

Assignment -2

Group 11

Aayush Amrute - amrute.a@northeastern.edu
(NUID : 002838262)

Paritosh Vyawahare - vyawahare.p@northeastern.edu
(NUID : 002416079)

Saurabh Chavan - chavan.sau@northeastern.edu
(NUID : 002479083)

Simran Sinha - sinha.sim@northeastern.edu
(NUID : 002475433)

Submission Date: 23-Oct-2024


Problem 1

Dataset- Activity

Dataset Description-Activity dataset shows how has global plastic waste disposal method changed over time.

TODO 1: Create an animated bar chart to illustrate how waste disposal methods have evolved over the years.





```
from google.colab import drive
drive.mount('/content/drive')
```

 Mounted at /content/drive

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib.animation as animation
```

```
# Here we will import the data first and then drop the null values
waste_data=pd.read_excel('/content/drive/MyDrive/Colab
Notebooks/Datasets/activity.xlsx')
cleaned_data = waste_data.dropna()
cleaned_data = cleaned_data[cleaned_data['Type'] != 'Generation']
```

cleaned_data

	year	Value	Type	
2	1960	5.6	Recycling	
3	1960	0.0	Combustion with energy recovery	
4	1960	82.5	Landfilling and other disposal	
7	1970	8.0	Recycling	
8	1970	0.5	Combustion with energy recovery	
9	1970	112.6	Landfilling and other disposal	
12	1980	14.5	Recycling	

13	1980	2.8	Combustion with energy recovery
14	1980	134.3	Landfilling and other disposal
18	1990	29.8	Combustion with energy recovery
19	1990	145.3	Landfilling and other disposal
17	1990	29.0	Recycling
16	1990	4.2	Composting*
21	2000	16.5	Composting*
22	2000	53.0	Recycling
23	2000	33.7	Combustion with energy recovery
24	2000	140.3	Landfilling and other disposal
29	2005	142.2	Landfilling and other disposal
28	2005	31.7	Combustion with energy recovery
27	2005	59.2	Recycling
26	2005	20.6	Composting*
31	2010	20.2	Composting*
32	2010	65.3	Recycling
33	2010	29.3	Combustion with energy recovery
34	2010	136.3	Landfilling and other disposal
38	2015	33.5	Combustion with energy recovery
39	2015	137.6	Landfilling and other disposal
36	2015	23.4	Composting*
37	2015	67.6	Recycling
41	2016	25.1	Composting*
			Recycling
42	2016	68.6	
43	2016	33.9	Combustion with energy recovery
44	2016	139.2	Landfilling and other disposal
46	2017	27.0	Composting*

47	2017	67.2	Recycling
48	2017	34.0	Combustion with energy recovery
49	2017	139.6	Landfilling and other disposal

Next steps:

[Generate code with cleaned_data](#)



[View recommended plots](#)

[New interactive sheet](#)

```
cleaned_data.sort_values(by='year', inplace=True)
figure, axis = plt.subplots(figsize=(10, 6))
```

```
def animate(current_year):
    axis.clear()
    data_for_current_year = cleaned_data[cleaned_data['year'] ==
current_year]
    sns.barplot(x='Value', y='Type', data=data_for_current_year,
palette='coolwarm', ax=axis)
    axis.set_title(f'Evolution of Plastic Waste Disposal
Methods in {current_year}', fontsize=12)
    axis.set_xlabel('Waste Amount (in tons)',
fontsize=12)
    axis.set_ylabel('Disposal Method', fontsize=12)
    axis.set_xlim(0, cleaned_data['Value'].max() * 1.1)
    for bar in axis.patches:
axis.text(bar.get_width() + 0.5, bar.get_y() + 0.55, f'{bar.get_width():.2f}', ha='c
```

```
all_years = sorted(cleaned_data['year'].unique())
animation_plot = animation.FuncAnimation(figure, animate, frames=all_years, repeat=False)
animation_plot.save('plastic_waste_animation.gif', writer='pillow', fps=1)
plt.show()
```

<ipython-input-62-bec2a27c8b39>:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.

<ipython-input-62-bec2a27c8b39>:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.

<ipython-input-62-bec2a27c8b39>:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.

<ipython-input-62-bec2a27c8b39>:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.

<ipython-input-62-bec2a27c8b39>:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.

<ipython-input-62-bec2a27c8b39>:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.

<ipython-input-62-bec2a27c8b39>:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.

<ipython-input-62-bec2a27c8b39>:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.

<ipython-input-62-bec2a27c8b39>:8: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.

<ipython-input-62-bec2a27c8b39>:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.

<ipython-input-62-bec2a27c8b39>:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.



TODO 2: What key insights can you draw from the visualization

- 1) **1960 Trends:** In 1960, land lling was the predominant method for disposing of plastic waste, with approximately 60 tons generated, indicating that it was the most widely used disposal method at the time. Only three methods were utilized in that year: land lling, recycling, and energy recovery.
- 2) **2000 Overview:** By the year 2000, the composting method was introduced as a waste disposal option. However, land lling continued to be the most commonly used method for disposing of plastic waste.
- 3) **2000-2010 Developments:** Between 2000 and 2010, there was a notable increase in the use of both combustion and recycling methods for plastic waste disposal, indicating a shift in waste management practices.
- 4) **Long-Term Trends:** Over the years, the reliance on land lling has declined, while the use of combustion and recycling methods has increased. This trend is a very positive development for

waste management, reflecting an improvement in sustainable practices and a move towards more environmentally friendly disposal methods.

🔗 Problem 2

Dataset - Global 500

Dataset Description - Fortune Global 500 List is a list of largest corporations worldwide which are measured by their total fiscal year revenues. Companies rankings sorted by total revenues for their respective fiscal years ended on or before March 31 of the relevant year.

TODO 1: Using treemap infer which countries dominate the global revenue landscape Interpret your key findings from the map

Hint: Perform data preprocessing before plotting the map

```
import plotly.express as px
global_data = pd.read_excel('/content/drive/MyDrive/Colab Notebooks/Datasets/Global 500.xlsx')


# We need to remove the '$' and unnecessary commas from revenue column and convert it to num
global_data['Revenues($millions)'] = global_data['Revenues($millions)'].replace({'\': '', ','})

# Aggregate total revenue by country
revenue_by_country = global_data.groupby('Country')['Revenues($millions)'].sum().reset_index # Sorting values
revenue_by_country = revenue_by_country.sort_values('Revenues($millions)', ascending=False)

# treemap
fig = px.treemap(revenue_by_country,
path=['Country'],
values='Revenues($millions)', title='Global Revenue Distribution by Country',
color='Revenues($millions)',
color_continuous_scale='Viridis')
fig.show()
```



Global Revenue Distribution by Country



Top Countries by Revenue: The plot clearly shows that the USA and China are the largest countries in terms of revenue, with the USA generating approximately \$8.48 trillion and China about 6.04 trillion. The USA is home to major companies like Amazon and Walmart, contributing significantly to its overall revenue.

Japan and Germany's Position: Following the USA and China, Japan and Germany rank third and fourth globally in revenue. Notable companies in these countries include Toyota and Mitsubishi, which play a crucial role in their economic output.

Similar Revenue Levels: Countries such as India, Brazil, Canada, and Taiwan have revenues that are nearly equal, each hovering around \$2 trillion. This indicates a competitive economic landscape among these nations.

Problem 3

Dataset - AirQualityUCI

Dataset Description - The dataset contains 9358 instances of hourly averaged responses from an array of 5 metal oxide chemical sensors embedded in an Air Quality Chemical Multisensor Device.

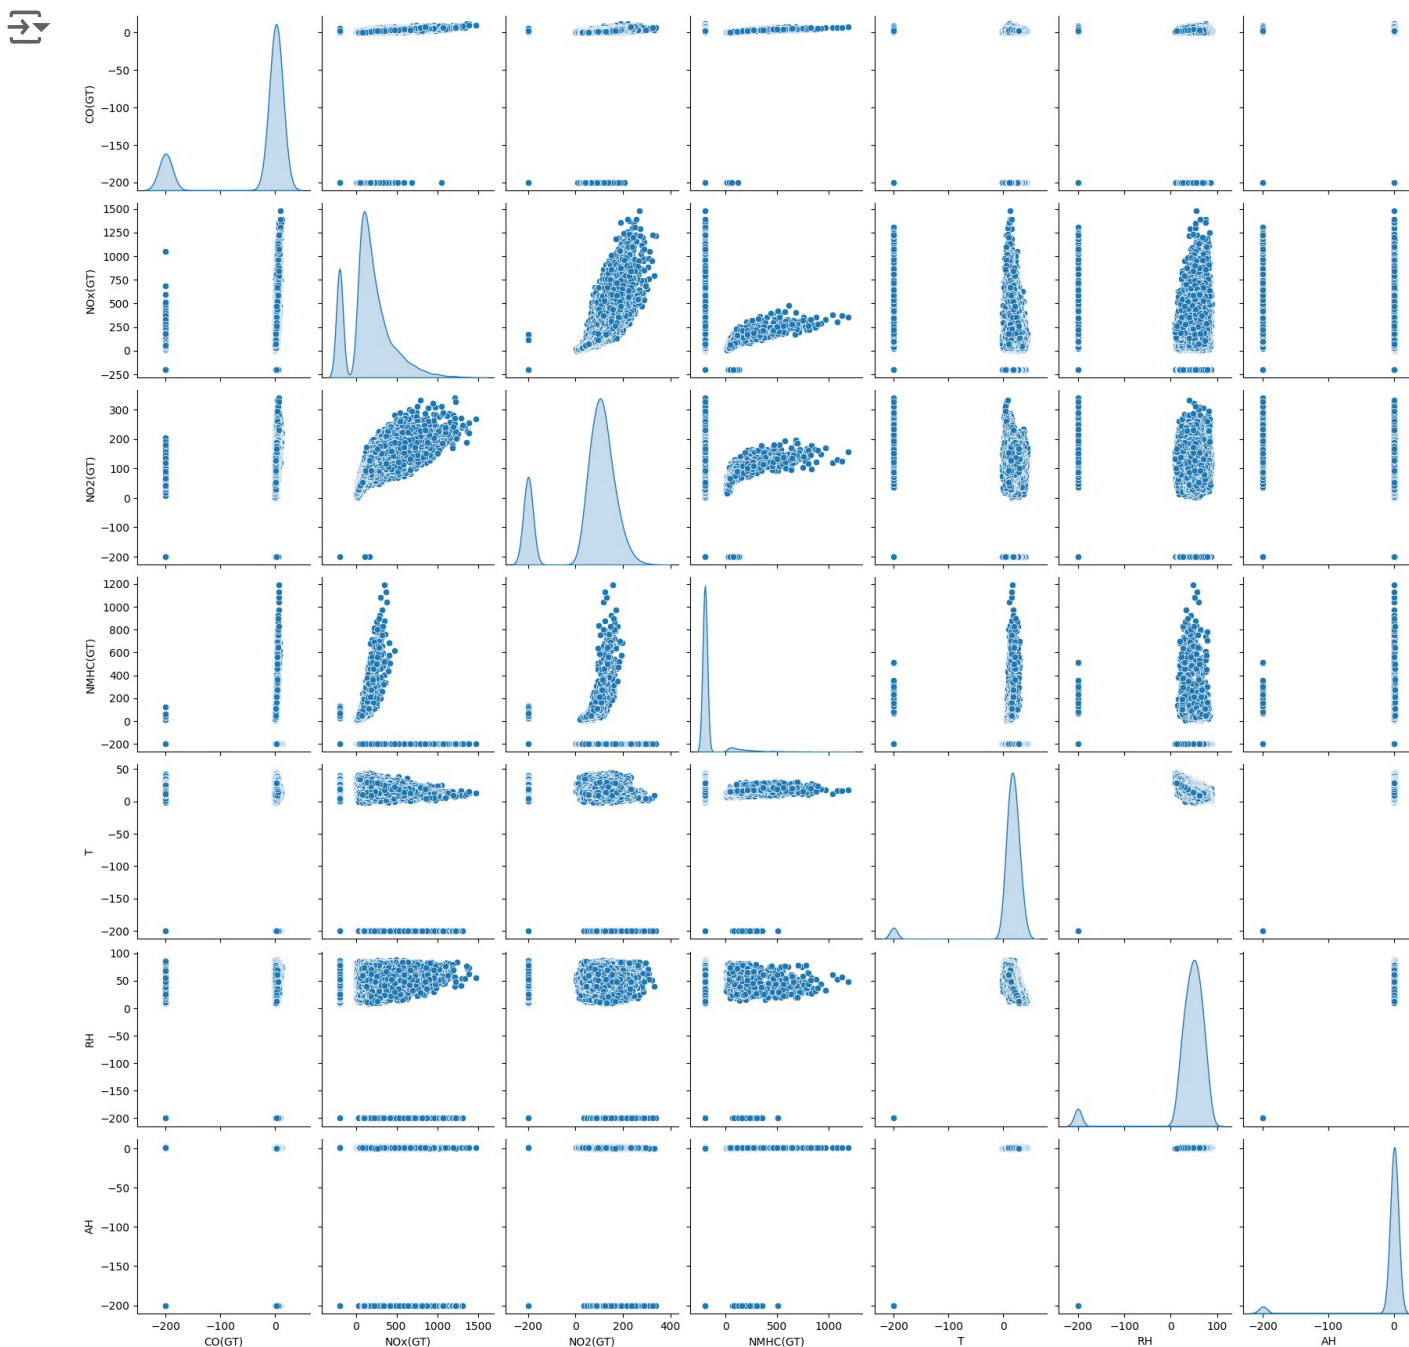
The device was located on the road in a significantly polluted area, at road level, within an Italian city. Data were recorded from March 2004 to February 2005 (one year) representing the longest freely available recordings of on-road deployed air quality chemical sensor devices responses. Ground Truth hourly averaged concentrations for CO, Non Methane Hydrocarbons, Benzene, Total Nitrogen Oxides (NOx) and Nitrogen Dioxide (NO2) and were provided by a co-located reference certified analyzer. Evidences of cross-sensitivities as well as both concept and sensor drifts are present as described in De Vito et al., Sens. And Act. B, Vol. 129,2,2008 (citation required) eventually affecting sensors concentration estimation capabilities. Missing values are tagged with -200 value.

TODO 1: Analyze the relationships between pollutants and environmental factors (T, RH, AH) using a scatter matrix (pair plot). Interpret your findings from the data

```
#Importing data
air_quality=pd.read_excel('/content/drive/MyDrive/Colab
Notebooks/Datasets/AirQualityUCI.xls')

#Taking only environmental factors which are needed
environmental_factor=['CO(GT)', 'NOx(GT)', 'NO2(GT)', 'NMHC(GT)', 'T', 'RH','AH']

sns.pairplot(air_quality[environmental_factor],diag_kind='kde')
plt.show()
```



Distribution of Pollutants: Pollutants such as CO(GT), NOx(GT), and NO2(GT) display right-skewed distributions, suggesting that lower concentrations are prevalent, while there are occasional instances of significantly higher levels.

Temperature Effects: There is a noticeable positive correlation between temperature (T) and pollutants like CO(GT), NMHC(GT), and NOx(GT). As temperatures rise, so do the levels of these pollutants, which could be attributed to heightened emissions from various human activities during warmer weather.

Humidity Relationships: Relative Humidity (RH) appears to have a slight negative correlation with certain pollutants, including CO(GT) and NOx(GT). This implies that higher humidity levels may be associated with lower concentrations of these pollutants, possibly because moisture in the air helps disperse them.

Humidity and Temperature Connection: Absolute Humidity (AH) exhibits a strong positive correlation with Temperature (T). This relationship is intuitive, as warmer air can hold more moisture, thereby directly influencing absolute humidity levels.

Problem 4

Dataset: [Wine Classification Dataset](#)

Dataset Description : These data are the results of a chemical analysis of wines grown in the same region in Italy but derived from three different cultivars. The analysis determined the quantities of 13 constituents found in each of the three types of wines.

TODO 1: Calculate the cumulative variance explained by each of the first two principal components using the raw data (without standardization). Explain how much of the total variance is captured by these two components.

```
# ucimlrepo package
!pip install ucimlrepo
```



Collecting ucimlrepo

```
Downloading ucimlrepo-0.0.7-py3-none-any.whl.metadata (5.5 kB)
Requirement already satisfied: pandas>=1.0.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: certifi>=2020.12.5 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: numpy>=1.22.4 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from
Downloading ucimlrepo-0.0.7-py3-none-any.whl (8.0 kB)
Installing collected packages: ucimlrepo
Successfully installed ucimlrepo-0.0.7
```

```
#Import the dataset into your code
from ucimlrepo import fetch_ucirepo
# fetch dataset wine =
fetch_ucirepo(id=109)
```

```
# data (as pandas
dataframes) X =
wine.data.features y =
wine.data.targets
```

```
wine.metadata
```

```
➔ {'uci_id': 109,  
  'name': 'Wine',  
  'repository_url': 'https://archive.ics.uci.edu/dataset/109/wine',  
  'data_url': 'https://archive.ics.uci.edu/static/public/109/data.csv',  
  'abstract': 'Using chemical analysis to determine the origin of wines',  
  'area': 'Physics and Chemistry',  
  'tasks': ['Classification'],  
  'characteristics': ['Tabular'],  
  'num_instances': 178,  
  'num_features': 13,  
  'feature_types': ['Integer', 'Real'],  
  'demographics': [],  
  'target_col': ['class'],  
  'index_col': None,  
  'has_missing_values': 'no',  
  'missing_values_symbol': None,  
  'year_of_dataset_creation': 1992,  
  'last_updated': 'Mon Aug 28 2023',  
  'dataset_doi': '10.24432/C5PC7J',  
  'creators': ['Stefan Aeberhard', 'M. Forina'],  
  'intro_paper': {'ID': 246,  
  
  'type': 'NATIVE',
```

```

'title': 'Comparative analysis of statistical pattern recognition methods in high
dimensional settings',
'authors': 'S. Aeberhard, D. Coomans, O. Vel',
'venue': 'Pattern Recognition',
'year': 1994,
'journal': None,
'DOI': '10.1016/0031-3203(94)90145-7',
'URL':
'https://www.semanticscholar.org/paper/83dc3e4030d7b9fbdbb4bde03ce12ab70ca10528',
'sha': None,
'corpus': None,
'arxiv': None,
'mag': None,
'acl': None,
'pmid': None,
'pmcid': None},
'additional_info': {'summary': 'These data are the results of a chemical analysis
of wines grown in the same region in Italy but derived from three different
cultivars. The analysis determined the quantities of 13 constituents found in each
of the three types of wines. \r\n\r\nI think that the initial data set had around 30
variables, but for some reason I only have the 13 dimensional version. I had a list
of what the 30 or so variables were, but a.) I lost it, and b.), I would not know
which 13 variables are included in the set.\r\n\r\nThe attributes are (dontated by
Riccardo Leardi, riclea@anchem.unige.it )\r\n1) Alcohol\r\n2) Malic acid\r\n3)
Ash\r\n4) Alcalinity of ash \r\n5) Magnesium\r\n6) Total phenols\r\n7)
Flavanoids\r\n8) Nonflavanoid phenols\r\n9) Proanthocyanins\r\n10)Color
intensity\r\n11)Hue\r\n12)OD280/OD315 of diluted wines\r\n13)Proline \r\n\r\nIn a
classification context, this is a well posed problem with "well behaved" class
structures. A good data set for first testing of a new classifier, but not very
challenging.
',
'purpose': 'test',
'funded_by': None,
'instances_represent': None,
'recommended_data_splits': None,
'sensitive_data': None,

```

wine.variables



	name	role	type	demographic	description	units	mis
0	class	Target	Categorical	None	None	None	
1	Alcohol	Feature	Continuous	None	None	None	
2	Malicacid	Feature	Continuous	None	None	None	
3	Ash	Feature	Continuous	None	None	None	
4	Alcalinity_of_ash	Feature	Continuous	None	None	None	
5	Magnesium	Feature	Integer	None	None	None	
6	Total_phenols	Feature	Continuous	None	None	None	
7	Flavanoids	Feature	Continuous	None	None	None	
8	Nonflavanoid_phenols	Feature	Continuous	None	None	None	
9	Proanthocyanins	Feature	Continuous	None	None	None	
10	Color_intensity	Feature	Continuous	None	None	None	
11	Hue	Feature	Continuous	None	None	None	
12	OD280_OD315_of_diluted_wines	Feature	Continuous	None	None	None	
13	OD280_OD315_of_diluted_wines	Feature	Continuous	None	None	None	

```
#PCA without standardizing
pca_raw = PCA(n_components=2)
pca_raw.fit(X)
```

```
explained_variance_raw = pca_raw.explained_variance_ratio_
cumulative_variance_raw = np.cumsum(explained_variance_raw)
```

```
print("Explained variance (raw data):", explained_variance_raw)
print("Cumulative variance (raw data):", cumulative_variance_raw)
```

```
➦ Explained variance (raw data): [0.99809123 0.00173592]
    Cumulative variance (raw data): [0.99809123 0.99982715]
```

TODO 2: Use the function PCA() on the centered but not scaled data to calculate the principal components. Compare how the components differ when using centered-only data versus raw data. Discuss any shifts in the proportion of variance explained.

```
#Finding centered values
```

```
centered_data = X -
np.mean(X,axis=0) pca_centered =
PCA(n_components=2)
pca_centered.fit(centered_data)
```

```
explained_variance_centered = pca_centered.explained_variance_ratio_
cumulative_variance_centered = np.cumsum(explained_variance_centered)
```

```
print("Explained variance (centered data):", explained_variance_centered)
print("Cumulative variance (centered data):", cumulative_variance_centered)
```

```
➦ Explained variance (centered data): [0.99809123 0.00173592]
    Cumulative variance (centered data): [0.99809123 0.99982715]
```

TODO 3: Use PCA on the standardized data but compute the top three principal components instead of the rst two. Compare how much more variance is captured by including the third principal component.

```
from sklearn.preprocessing import StandardScaler
```

```
#Standardizing Data scaler =
StandardScaler() standardized_data =
scaler.fit_transform(X)
```

```
pca_standardized = PCA(n_components=3)
pca_standardized.fit(standardized_data)
```

```
explained_variance_standardized = pca_standardized.explained_variance_ratio_
cumulative_variance_standardized = np.cumsum(explained_variance_standardized)
```

```
print("Explained variance (standardized data):", explained_variance_standardized)
print("Cumulative variance (standardized data):",
      cumulative_variance_standardized)
```

```
➡ Explained variance (standardized data): [0.36198848 0.1920749 0.11123631]
   Cumulative variance (standardized data): [0.36198848 0.55406338 0.66529969]
```

TODO 4: Compare the results of PCA on standardized vs. min-max normalized data. Discuss the impact of these two techniques on the PCA outcomes.

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

```
# z-score normalization scaler_standard =
StandardScaler() standardized_data =
scaler_standard.fit_transform(X)
```

```
pca_standardized = PCA(n_components=3)
pca_standardized.fit(standardized_data)
```

```
explained_variance_standardized = pca_standardized.explained_variance_ratio_
cumulative_variance_standardized = np.cumsum(explained_variance_standardized)
```

```
print("Explained variance (standardized data):", explained_variance_standardized)
print("Cumulative variance (standardized data):",
      cumulative_variance_standardized)
```

```
# Normalizing data scaler_minmax =
MinMaxScaler() normalized_data =
scaler_minmax.fit_transform(X)
```

```
pca_normalized = PCA(n_components=3)
pca_normalized.fit(normalized_data)
```

```
explained_variance_normalized = pca_normalized.explained_variance_ratio_
cumulative_variance_normalized = np.cumsum(explained_variance_normalized)
```

```
print("Explained variance (min-max normalized data):", explained_variance_normalized)
print("Cumulative variance (min-max normalized data):", cumulative_variance_normalized)
```


```
➡ Explained variance (standardized data): [0.36198848 0.1920749 0.11123631]
   Cumulative variance (standardized data): [0.36198848 0.55406338 0.66529969]
   Explained variance (min-max normalized data): [0.40749485 0.18970352 0.08561671]
   Cumulative variance (min-max normalized data): [0.40749485 0.59719836 0.68281507]
```


TODO 5: Instead of plotting only the rst two PCs, generate a 3D plot using the rst three principal components. Use color to differentiate wine classes and interpret the additional insights from the third component. python

```
from sklearn.decomposition import PCA
from mpl_toolkits.mplot3d import Axes3D
```

```
X_pca_std = pca_standardized.fit_transform(standardized_data)
pc1 = X_pca_std[:, 0] pc2 = X_pca_std[:, 1] pc3 =
X_pca_std[:, 2]
```

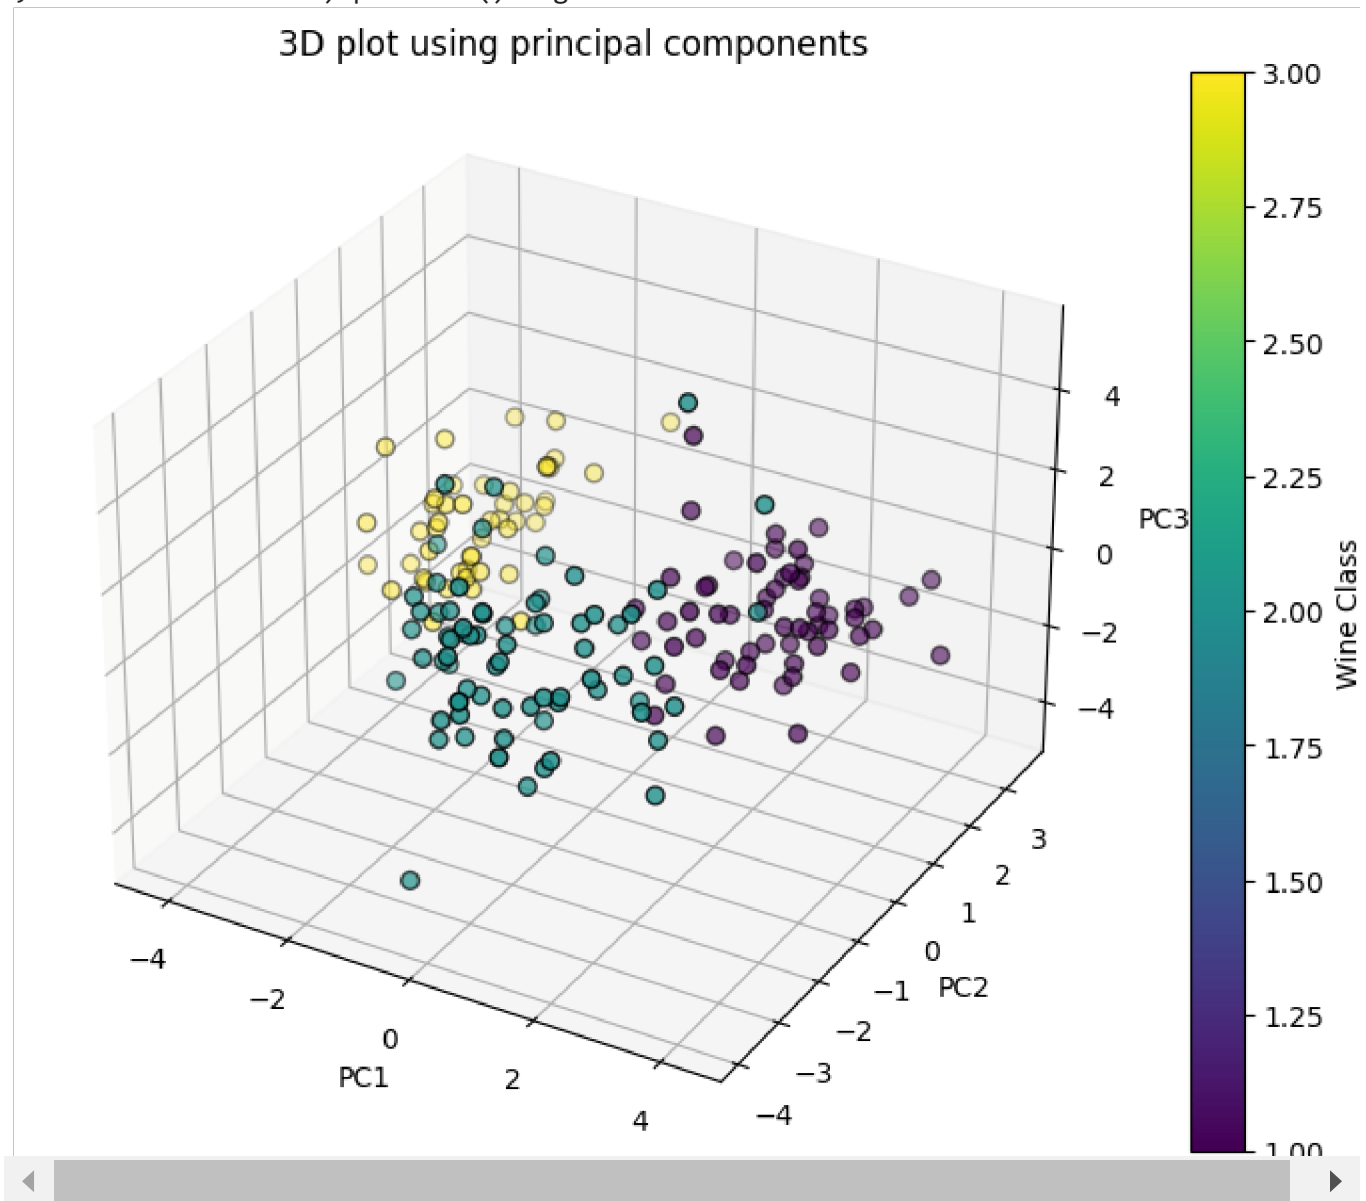
```
fig = plt.figure(figsize=(10, 7)) ax =
fig.add_subplot(111, projection='3d')
➡
```



```
scatter = ax.scatter(pc1, pc2, pc3, c=y, cmap='viridis', edgecolor='k', s=40)
```

```
#plotting the data
```

```
ax.set_title("3D plot using principal components")
ax.set_xlabel("PC1") ax.set_ylabel("PC2")
ax.set_zlabel("PC3") color = fig.colorbar(scatter,
ax=ax, label='Wine Class') plt.show() fig
```



Principal Component 1 (PC1): Represents the highest amount of variance in the dataset, highlighting the most pronounced differences among the wine samples.

Principal Component 2 (PC2): Accounts for the second largest variance, capturing variations that are independent of those represented by PC1.

Principal Component 3 (PC3): Captures additional variance that may not be explained by PC1 and PC2, allowing for more detailed differentiation of the samples.

Separation of Wine Classes: The analysis reveals a distinct separation among the three wine categories, illustrating the effectiveness of PCA in distinguishing between them.

Class Distinction: PC1 is particularly effective in differentiating between the various wine classes.

Further Differentiation: Both PC2 and PC3 contribute to refining the separation of overlapping classes, with PC3 revealing subtle distinctions in a three-dimensional representation.

Problem 5

Dataset: Life Expectancy

Introduction: The above dataset gives life expectancy related data for 37 countries in 2014.

Consider only the following variables in your analysis: 'GDP', 'Income composition of resources', 'Schooling', and 'Total expenditure'.

TODO 1: Perform Z-score normalization on the numeric variables to scale the data. Compare the distribution of features before and after Z-score normalization and discuss the effect on variance.

```
#Importing Data df=pd.read_csv('/content/drive/MyDrive/Colab Notebooks/Datasets/Life Expectancy.csv')
```

```
df.head()
```



	Country	Year	Status	Life expectancy	Adult Mortality	infant deaths	Alcohol	percentage Hep expenditure
0	Afghanistan	2014	Developing	59.9	271	64	0.01	73.523582
1	Australia	2014	Developed	82.7	6	1	9.71	10769.363050
2	Austria	2014	Developed	81.4	66	0	12.32	8350.193523
3	Bangladesh	2014	Developing	71.4	132	98	0.01	10.446403
4	Belgium	2014	Developed	89.0	76	0	12.60	7163.348923

5 rows × 22 columns

```
numeric_cols = df.select_dtypes(include=['float64', 'int64']).columns
```

```
df_numeric = df[numeric_cols] scaler = StandardScaler() df_numeric_normalized = pd.DataFrame(scaler.fit_transform(df_numeric), columns=numeric_cols) print("Z-score Normalized Data:") print(df_numeric_normalized.head())
```



Z-score Normalized Data:

	Year	Life expectancy	Adult Mortality	infant deaths	Alcohol
0	0.0	-1.517367	1.243871	-0.046452	-1.082462
1	0.0	0.911111	-1.242096	-0.401464	1.159808
2	0.0	0.772645	-0.679236	-0.407099	1.763140

3	0.0	-0.292477	-0.060089	0.145142	-1.082462	4	0.0
	1.582138	-0.585425	-0.407099	1.827866			

	percentage expenditure	Hepatitis B	Measles	BMI	\
0	-0.526543	-0.538579	-0.317693	-1.158468	
1	2.048225	0.577049	-0.326338	1.155703	
2	1.465868	0.846338	-0.339019	0.717228	
3	-0.541727	0.807868	-0.329238	-1.202315	4
	1.180163	0.846338	-0.341692	1.024160	

	under-five deaths	Polio	Total expenditure	Diphtheria	HIV/AIDS	\
0	-0.035392	-0.970515	0.619128	-0.891875	-0.411981	
1	-0.401541	0.437683	1.017244	0.436959	-0.411981	
2	-0.405849	0.686189	1.591945	0.702725	-0.411981	
3	0.115375	0.644771	-1.101762	0.658431	-0.411981	
4	-0.401541	0.727607	-1.496667	0.747020	-0.411981	

	GDP	Population	thinness	1-19 years	thinness 5-9 years	\
0	-0.643117	-0.297465		1.971357	1.801683	
1	2.180926	-0.287950		-0.692190	-0.781652	
2	1.681597	-0.258756		-0.503062	-0.567648	
3	-0.662744	-0.223865		2.065921	1.969830	
4	1.503576	-0.298476		-0.629147	-0.720508	

	Income composition of resources	Schooling
0	-1.567660	-1.158371
1	1.236204	2.140971
2	0.968009	0.713371
3	-0.994696	-1.158371
4	0.955818	0.840269

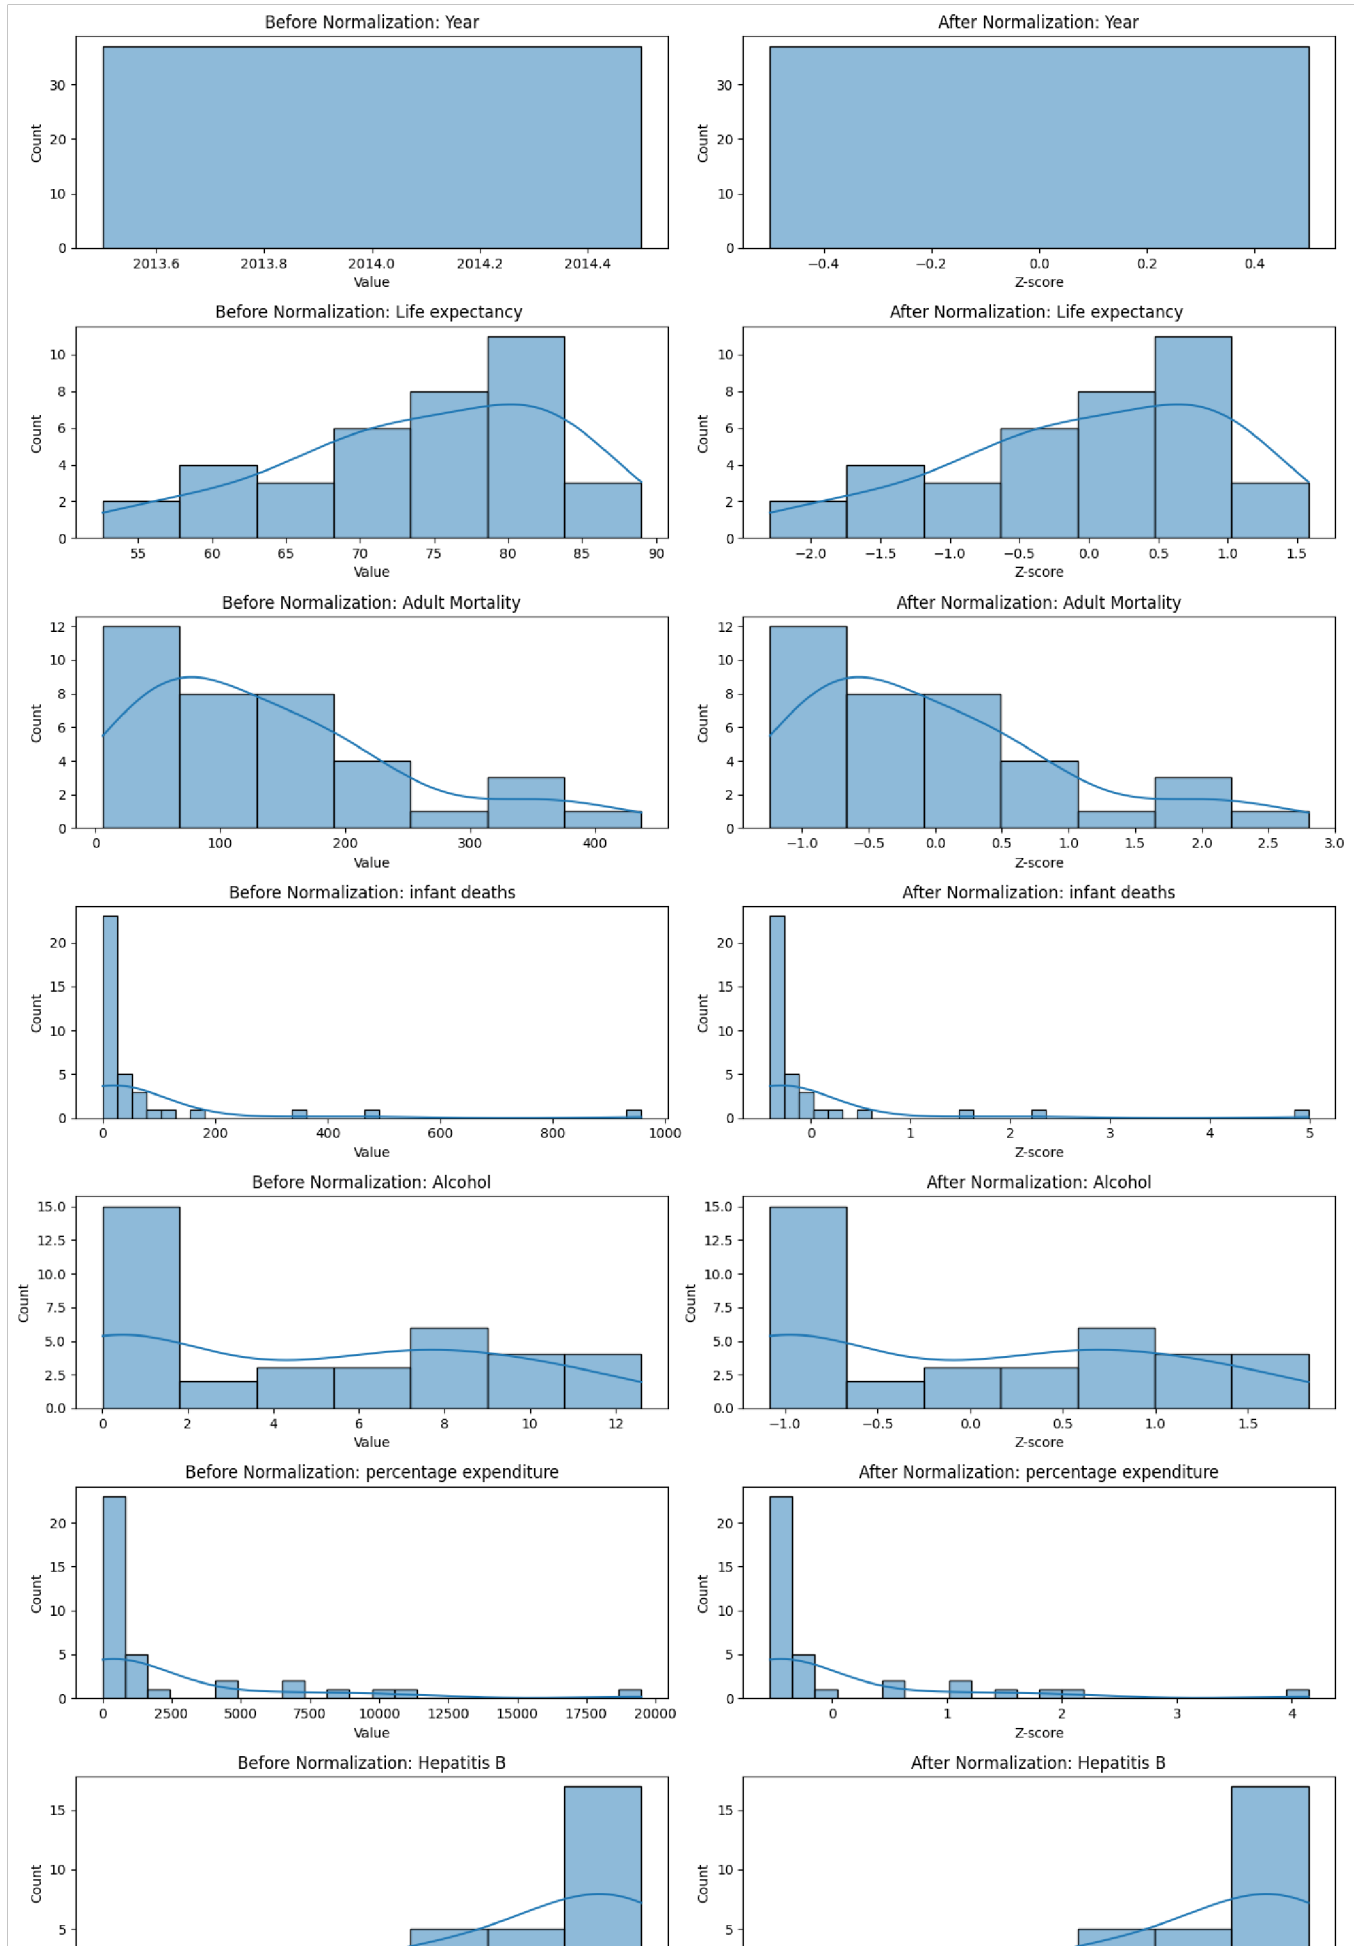
```
import matplotlib.pyplot as plt
import seaborn as sns
```

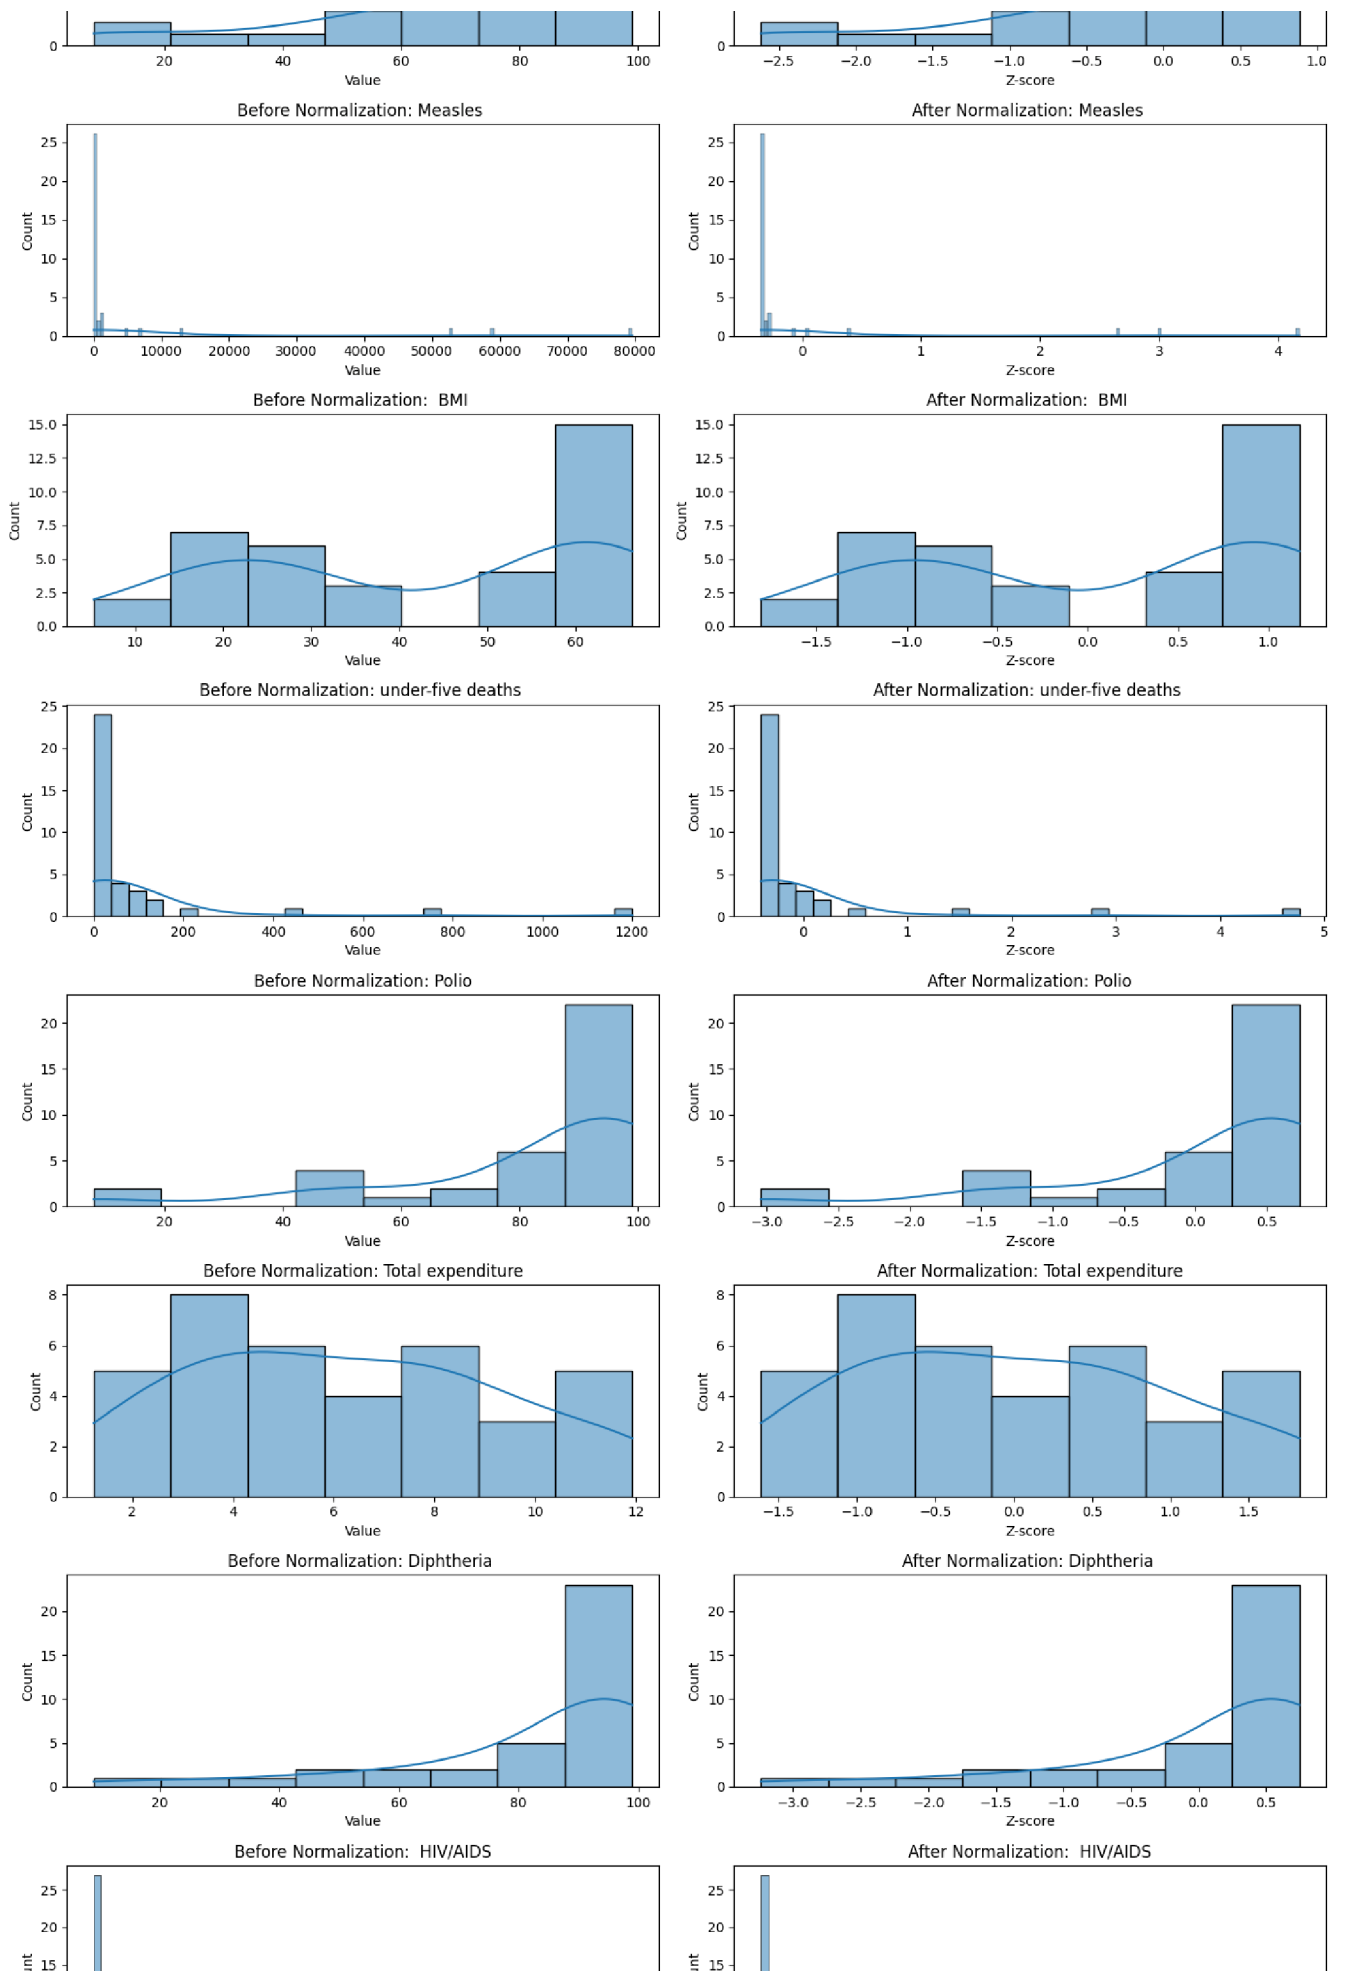
```
# Create a copy for comparison df_before_norm = df_numeric.copy()
n_cols = 2 n_rows = len(numeric_cols) fig, axes =
plt.subplots(n_rows, n_cols, figsize=(14, n_rows * 3))
```

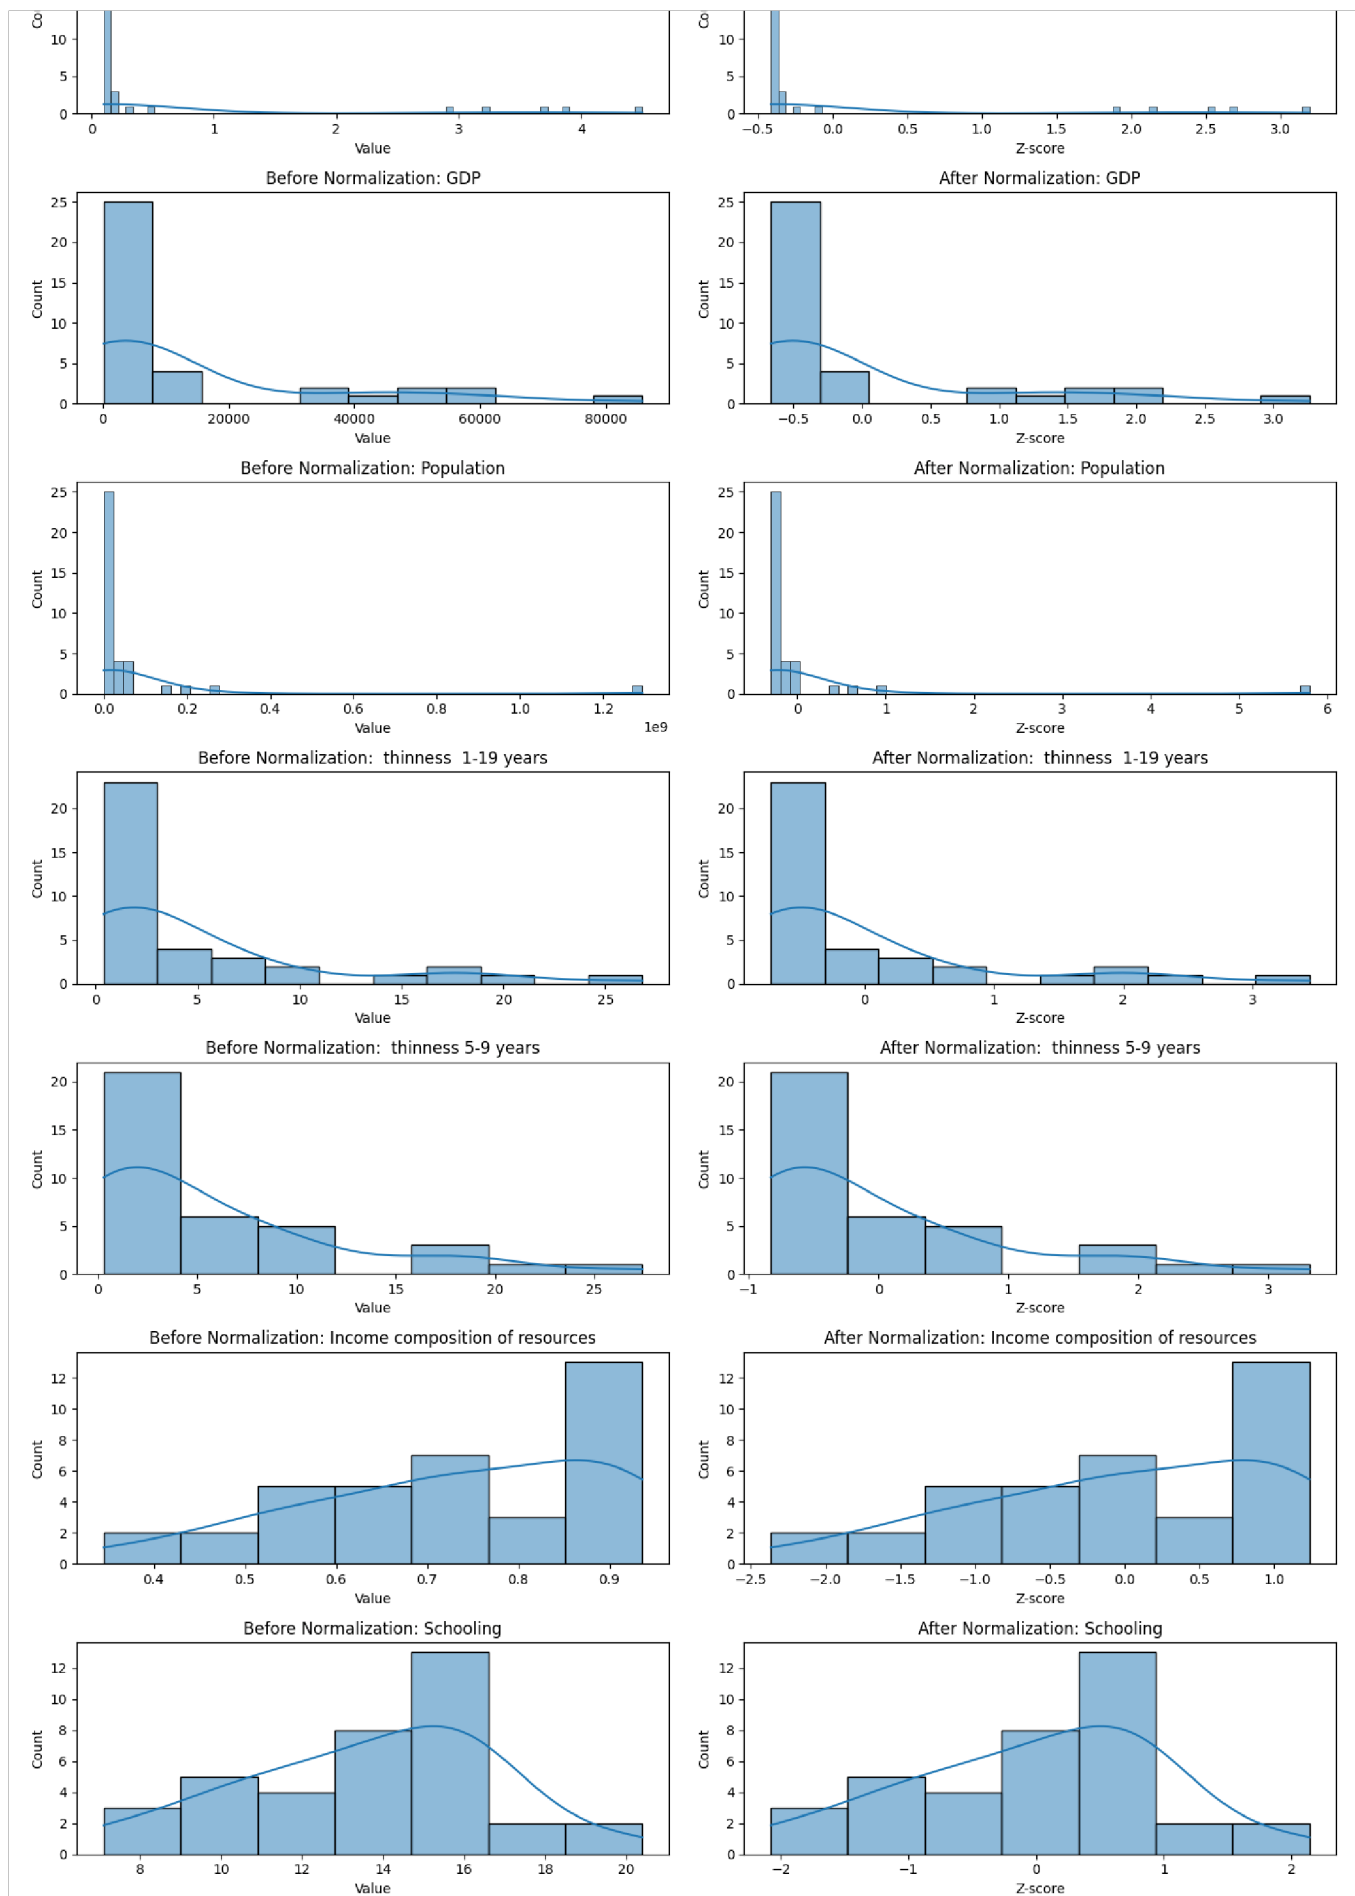
```
for i, col in enumerate(numeric_cols):
    sns.histplot(df_before_norm[col], ax=axes[i, 0], kde=True)
    axes[i, 0].set_title(f'Before Normalization: {col}')
    axes[i, 0].set_xlabel('Value')
    sns.histplot(df_numeric_normalized[col], ax=axes[i, 1],
kde=True) axes[i, 1].set_title(f'After Normalization: {col}')
    axes[i, 1].set_xlabel('Z-score')
plt.tight_layout()
plt.show()
```

```
print("Variance before normalization:")
print(df_before_norm.var())
```

```
print("\nVariance after normalization (should be close to 1):")  
print(df_numeric_normalized.var())
```







Variance before normalization:

Year	0.000000e+00	Life expectancy	9.059422e+01
Adult Mortality	1.167886e+04	infant deaths	3.236630e+04
Alcohol	1.923388e+01	percentage expenditure	1.773588e+07
Hepatitis B	6.961818e+02		
Measles	3.177893e+08		
BMI	4.330079e+02		
under-five deaths	5.538873e+04		
Polio	5.991411e+02		
Total expenditure	9.970631e+00		
Diphtheria	5.238423e+02		
HIV/AIDS	1.530270e+00		
GDP	4.890421e+08		
Population	4.627788e+16		
thinness 1-19 years	4.137632e+01		
thinness 5-9 years	4.398565e+01		
Income composition of resources	2.766310e-02		
Schooling	1.021201e+01		

dtype: float64

Variance after normalization (should be close to 1):

Year	0.000000
Life expectancy	1.027778
Adult Mortality	1.027778
infant deaths	1.027778
Alcohol	1.027778
percentage expenditure	1.027778
Hepatitis B	1.030303
Measles	1.027778
BMI	1.027778
under-five deaths	1.027778
Polio	1.027778
Total expenditure	1.027778
Diphtheria	1.027778
HIV/AIDS	1.027778
GDP	1.027778
Population	1.027778
thinness 1-19 years	1.027778
thinness 5-9 years	1.027778
Income composition of resources	1.027778
Schooling	1.027778

dtype: float64

🔍 Z-Score Normalization and Its Effects on Variance

In this analysis, we performed **Z-score normalization** on all numeric variables in the Life Expectancy dataset. The goal of Z-score normalization is to standardize the data so that each feature has a mean of 0 and a variance of 1. This process allows us to scale features that were originally on different scales and variances, ensuring that no single feature dominates the analysis due to its scale.

Before Normalization:

The dataset showed significant variation in the scales of different features. For instance, **GDP** had a variance of **489,042,100**, while **Schooling** had a variance of only **10.21**. Such disparities suggest that the larger variance features (like **GDP** or **Population**) would heavily influence any statistical model or analysis, potentially leading to biased results.

After Normalization:

After applying **Z-score normalization**, the variance of every feature was approximately **1**. This means that all features now contribute equally in terms of their variability. Features with larger scales, like **GDP** or **Population**, no longer overpower features with smaller scales like **Schooling** or **Alcohol**.


Conclusion:

Z-score normalization effectively standardizes the dataset, ensuring that each feature has equal weight in the analysis. This is crucial for models that are sensitive to feature scaling, as it prevents certain features from dominating due to their original measurement units or variances.

As a result, this step ensures a **fair comparison** between features and is an important preprocessing step for improving **model accuracy** and **interpretability**.

TODO 2: Covariance Matrix and Eigen Decomposition (PCA) This step involves performing Principal Component Analysis (PCA) for dimensionality reduction. You'll compute the covariance matrix and use eigen decomposition to find principal components. `import numpy as np`

```
df_numeric_filled = df_numeric.copy() df_numeric_filled[numeric_cols] =  
df_numeric[numeric_cols].fillna(df_numeric[numeric_cols].mean()) cov_matrix =  
np.cov(df_numeric_filled[numeric_cols].values.T) print("Covariance Matrix:\n", cov_matrix)
```

 4.33007853e+02 -1.80325075e+03 2.00815165e+02 2.40947583e+01
1.98708559e+02 -1.25132733e+01 2.30915741e+05 -1.06166176e+09 -
7.32665691e+01 -9.00499775e+01 2.62613198e+00 4.65111411e+01]
[0.00000000e+00 -9.08668544e+02 7.52349324e+03 4.21408626e+04
-2.71811712e+02 -2.03191396e+05 -5.31416667e+02 2.68754745e+06

```
[ 0.00000000e+00  1.48216884e+00 -1.37235233e+01 -1.04882695e+01
 5.09668086e-01  3.77915738e+02  1.50252778e+00 -4.56036791e+02
 2.62613198e+00 -1.42910420e+01  2.11619369e+00  1.90357740e-01
 2.07483483e+00 -1.28022748e-01  2.18514181e+03 -5.56137952e+06
-6.33703979e-01 -7.55105405e-01  2.76631021e-02  4.97145571e-01]
[ 0.00000000e+00  2.67222973e+01 -2.49982508e+02 -2.05957282e+02
 9.90652402e+00  7.28781678e+03  2.67055556e+01 -9.80454077e+03
 4.65111411e+01 -2.77105856e+02  3.64327327e+01  3.97888664e+00
 3.57595345e+01 -2.08569820e+00  4.26293777e+04 -1.00249753e+08
-1.21140165e+01 -1.43221021e+01  4.97145571e-01  1.02120120e+01]]
```

TODO 3: Variance Explained by Each Principal Component This step involves calculating how much variance each principal component explains.

```
eigenvalues, eigenvectors = np.linalg.eig(cov_matrix)

print("\nEigenvalues:\n", eigenvalues)
print("\nEigenvectors:\n", eigenvectors)

sorted_indices = np.argsort(eigenvalues)[::-1] # Sort in descending order
sorted_eigenvalues = eigenvalues[sorted_indices]
sorted_eigenvectors = eigenvectors[:, sorted_indices]

print("\nSorted Eigenvalues:\n", sorted_eigenvalues)
print("\nSorted Eigenvectors:\n", sorted_eigenvectors)
```



```

1.02692962e-01  5.26680475e-02 -2.51510452e-03  2.16148284e-01
1.40854750e-01  9.56314387e-01  4.32371900e-02  0.00000000e+00] [-
1.38788452e-05  9.78531329e-01  9.59132447e-02 -1.82261356e-01
5.25842295e-03  5.17929844e-03 -1.58832046e-03 -8.46815938e-04
3.05430760e-04  1.56770982e-04  1.24238512e-04  8.19569723e-05
9.28291143e-06 -6.22241511e-05 -1.35593715e-04 -3.46286336e-05
7.52048781e-05 -3.04829392e-05 -2.32112543e-06  0.00000000e+00] [-
9.99999998e-01  1.93825755e-05 -5.42276669e-05 -6.42325891e-07
-9.09363513e-07  2.44011421e-07 -2.03792944e-09  5.86180670e-09
6.07056929e-09 -5.93683057e-08 -4.78823493e-08  1.42408197e-08 -
5.36176179e-09  4.47670129e-09  6.52630360e-10 -1.72095896e-09
5.46391728e-09 -1.31696236e-09 -8.24247600e-11  0.00000000e+00] [-
1.74735559e-08 -7.96207523e-05 -6.41863048e-05  4.97051747e-04
6.80726508e-03  8.09031109e-03  3.99476087e-02 -1.57359552e-01 -
5.36465872e-02  3.40280559e-01 -7.55128931e-02  1.99584781e-01
-4.50409802e-01 -3.29287535e-01  1.14012396e-01  6.88670783e-01
4.47100369e-02 -9.35612146e-02 -2.59819520e-03  0.00000000e+00] [-
1.69767823e-08 -1.02740967e-04  5.29427858e-06  8.00759789e-04
1.30252501e-02  1.08782171e-02  4.21893852e-02 -1.64542258e-01 -
5.27998888e-02  2.47977926e-01 -1.35659733e-01  1.97119062e-01
-4.84449719e-01 -3.19848122e-01  1.32573278e-02 -6.77150314e-01
-4.45345465e-02  2.27010771e-01  9.84842844e-03  0.00000000e+00]
[ 1.20173603e-10  4.28851535e-06  3.57819859e-07 -3.09954686e-05
-3.59639870e-04 -7.95460957e-04  3.04115001e-04  3.13636223e-03
3.11323734e-03 -3.44911764e-03  3.37303034e-03 -7.67789173e-03 -
5.51382593e-03 -1.86395910e-03 -6.08903568e-03  4.04106678e-04
-1.05307640e-02 -4.27850330e-02  9.98941850e-01  0.00000000e+00]
[ 2.16625600e-09  8.41269342e-05 -9.48582249e-07 -7.12535131e-04
-6.93318308e-03 -1.29028486e-02  1.30302989e-03  4.91811768e-02
5.08961198e-02 -9.47519258e-02  2.99295592e-02 -1.53693154e-01 -
7.31950392e-02 -1.11428700e-01 -1.17781462e-01  1.07990085e-01
-9.48786040e-01  1.33107880e-01 -7.61025589e-03  0.00000000e+00]]

```

```
explained_variance_ratio = eigenvalues/np.sum(eigenvalues) print("Explained
```

```
Variance by each Principal Component:", explained_variance_ratio)
```

```

➡ Explained Variance by each Principal Component: [9.99999985e-01 1.07990784e-08 3.6833716
5.90284440e-13 1.77075055e-13 2.21080425e-14 5.33867498e-15 3.80609095e-
15 1.51743076e-15 1.12772442e-15 3.69438102e-16
2.53712985e-16 1.62128802e-16 8.27702013e-17 3.36695749e-17
2.31379426e-17 5.59021191e-18 9.29203283e-21 0.00000000e+00]

```

TODO 4: Reconstructing the Data Using Principal Components This task would involve using the principal components to approximate the original data by reconstructing it from the top few principal components.

```
k = 5 top_k_components =
sorted_eigenvectors[:, :k]
```

```
projected_data = np.dot(df_numeric_normalized, top_k_components)
```

```
reconstructed_data = np.dot(projected_data, top_k_components.T)
```

```
reconstructed_df = pd.DataFrame(reconstructed_data, columns=numeric_cols)
```

```
print("Reconstructed Data (using top 5 principal components):")
print(reconstructed_df.head())
```



Reconstructed Data (using top 5 principal components):

	Year	Life expectancy	Adult Mortality	infant deaths	Alcohol	\
0	0.0	-0.011230	0.106865	0.240210	-0.003006	
1	0.0	0.024098	-0.225267	-0.536373	0.006616	
2	0.0	0.020992	-0.197572	-0.460648	0.005702	
3	0.0	-0.004820	0.044533	0.108837	-0.001365	4 0.0
	0.021196	-0.200790	-0.458546	0.005700		

	percentage expenditure	Hepatitis B	Measles	BMI	\
0	-0.483548	-0.007132	-0.318708	-0.010868	
1	2.014013	0.009474	-0.325682	0.021145	
2	1.439908	0.010120	-0.338405	0.019218	
3	-0.548349	-0.001260	-0.328882	-0.003885	4
	1.150983	0.012050	-0.341098	0.020176	

	under-five deaths	Polio	Total expenditure	Diphtheria	HIV/AIDS	\
0	0.356724	-0.009975	-0.001941	-0.012211	0.001519	
1	-0.796178	0.018773	0.004853	0.023441	-0.003286	
2	-0.683713	0.017161	0.004004	0.021311	-0.002849	
3	0.161736	-0.003532	-0.001055	-0.004387	0.000665	
4	-0.680554	0.018127	0.003825	0.022391	-0.002864	

	GDP	Population	thinness	1-19 years	thinness 5-9 years	\
0	-0.654970	-0.297465		0.002969	0.005650	
1	2.190484	-0.287949		-0.006301	-0.012178	
2	1.688719	-0.258755		-0.005490	-0.010592	
3	-0.661805	-0.223865		0.001287	0.002439	
4	1.511349	-0.298476		-0.005548	-0.010679	

	Income composition of resources	Schooling
0	-0.000154	-0.002929
1	0.000326	0.006110
2	0.000286	0.005383
3	-0.000065	-0.001201
4	0.000290	0.005495

TODO 5: Compare Original Data with Reconstructed Data The task here is to compare the original data with the reconstructed data to see the difference between the two.

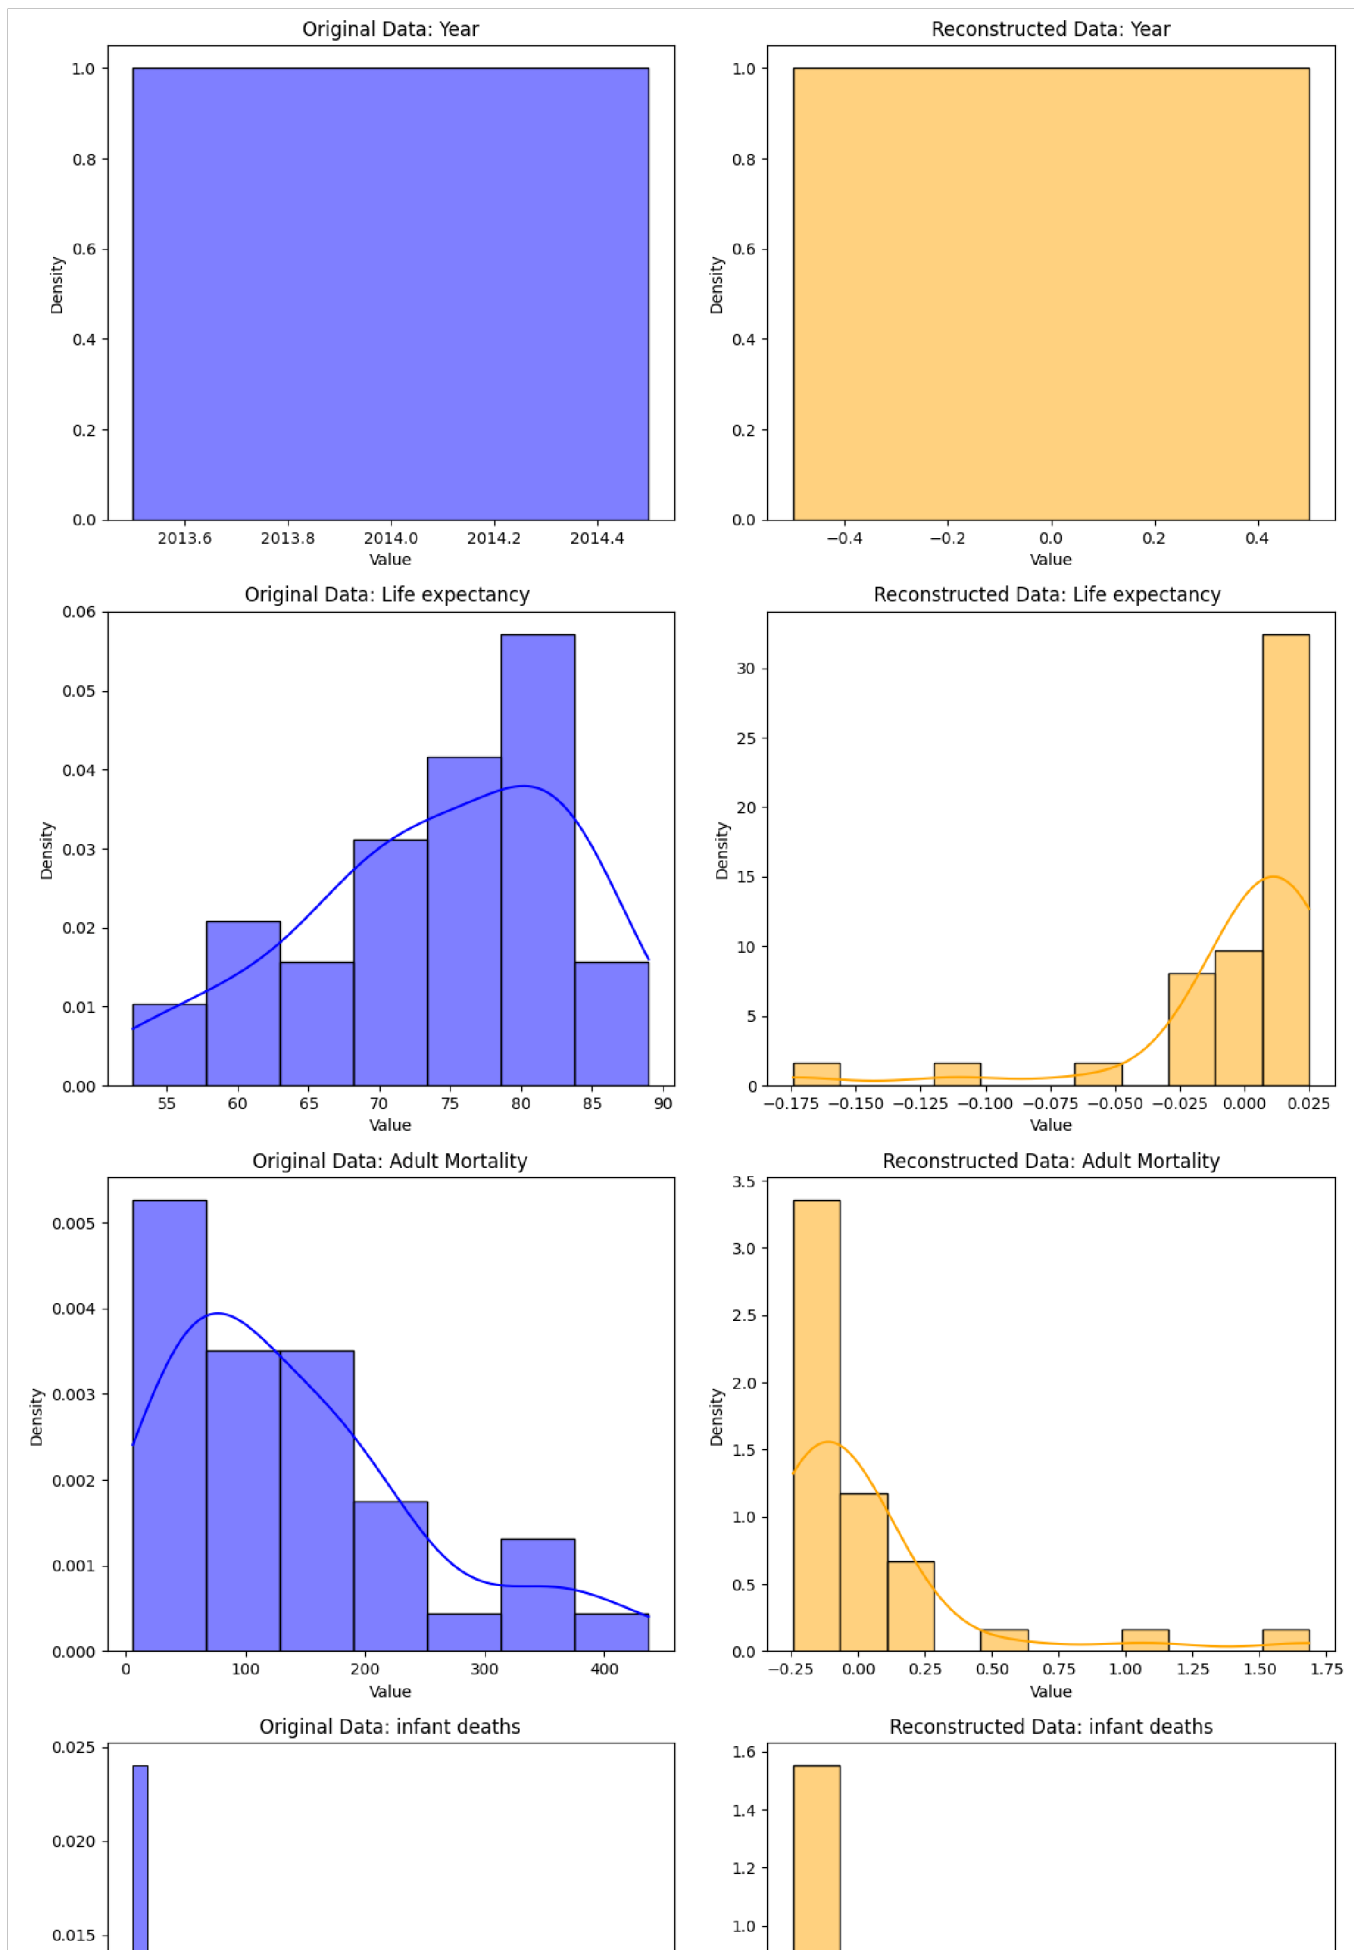
```
columns_to_compare = numeric_cols


n_cols = len(columns_to_compare) fig, axes =
plt.subplots(n_cols, 2, figsize=(12, 5 * n_cols))

for i, col in enumerate(columns_to_compare):
    sns.histplot(df_numeric_filled[col], ax=axes[i, 0], kde=True, color='blue', label='Original Data')
    sns.histplot(reconstructed_df[col], ax=axes[i, 1], kde=True, color='orange', label='Reconstructed Data')
    axes[i, 0].set_title(f'Original Data: {col}') axes[i, 0].set_xlabel('Value')
    axes[i, 0].set_ylabel('Density') axes[i, 1].set_title(f'Reconstructed Data: {col}')
    axes[i, 1].set_xlabel('Value') axes[i, 1].set_ylabel('Density')


plt.tight_layout()
plt.show()

difference = df_numeric_filled[columns_to_compare] - reconstructed_df
print("\nDifference between Original and Reconstructed Data (first 5 rows):")
print(difference.head())
```







Original Data: Alcohol



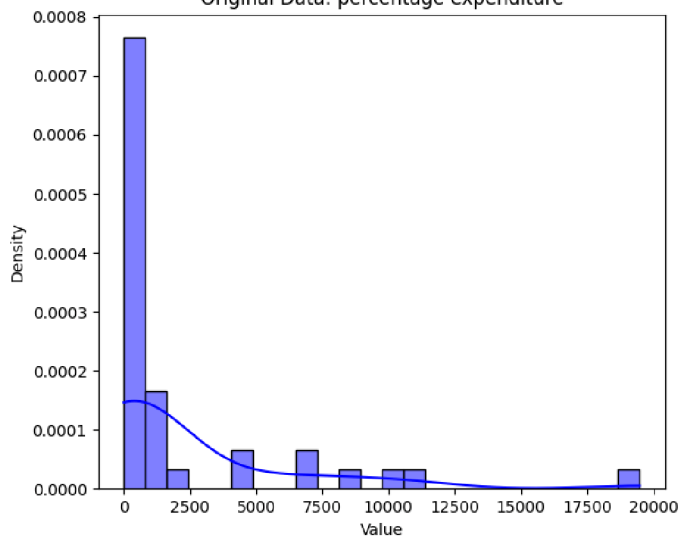
Reconstructed Data: Alcohol



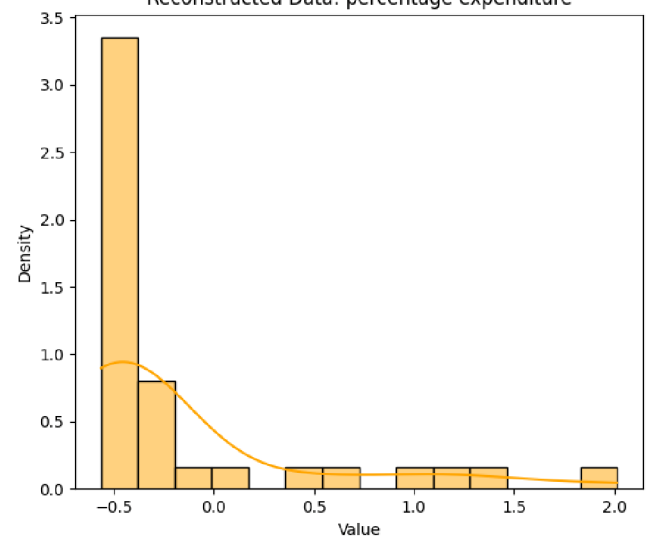
Original Data: percentage expenditure




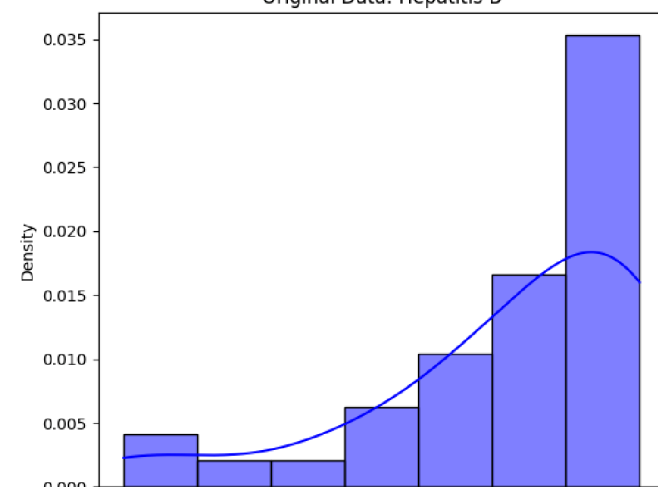
Reconstructed Data: percentage expenditure



Original Data: Hepatitis B



Reconstructed Data: Hepatitis B



Difference between Original and Reconstructed Data (first 5 rows):

	Year	Life expectancy	Adult Mortality	infant deaths	Alcohol \
0	2014.0	59.911230	270.893135	63.759790	0.013006
1	2014.0	82.675902	6.225267	1.536373	9.703384
2	2014.0	81.379008	66.197572	0.460648	12.314298
3	2014.0	71.404820	131.955467	97.891163	0.011365
4	2014.0	88.978804	76.200790	0.458546	12.594300

	percentage expenditure	Hepatitis B	Measles	BMI \
0	74.007130	62.007132	492.318708	18.610868
1	10767.349037	90.990526	340.325682	66.078855
2	8348.753615	97.989880	117.338405	57.080782
3	10.994752	97.001260	289.328882	17.703885
4	7162.197940	97.987950	70.341098	63.379824

	under-five deaths	Polio	Total expenditure	Diphtheria	HIV/AIDS \
0	85.643276	58.009975	8.181941	62.012211	0.098481
1	1.796178	91.981227	9.415147	91.976559	0.103286
2	0.683713	97.982839	11.205996	97.978689	0.102849
3	120.838264	97.003532	2.821055	97.004387	0.099335
4	1.680554	98.981873	1.586175	98.977609	0.102864

	GDP	Population	thinness	1-19 years	thinness 5-9 years \
0	613.351484	3.275823e+05		17.497031	17.494350

...

1		0.935674	20.393890
2		0.891714	15.894617
3		0.570065	10.001201
4		0.889710	16.294505

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output [settings](#)...